

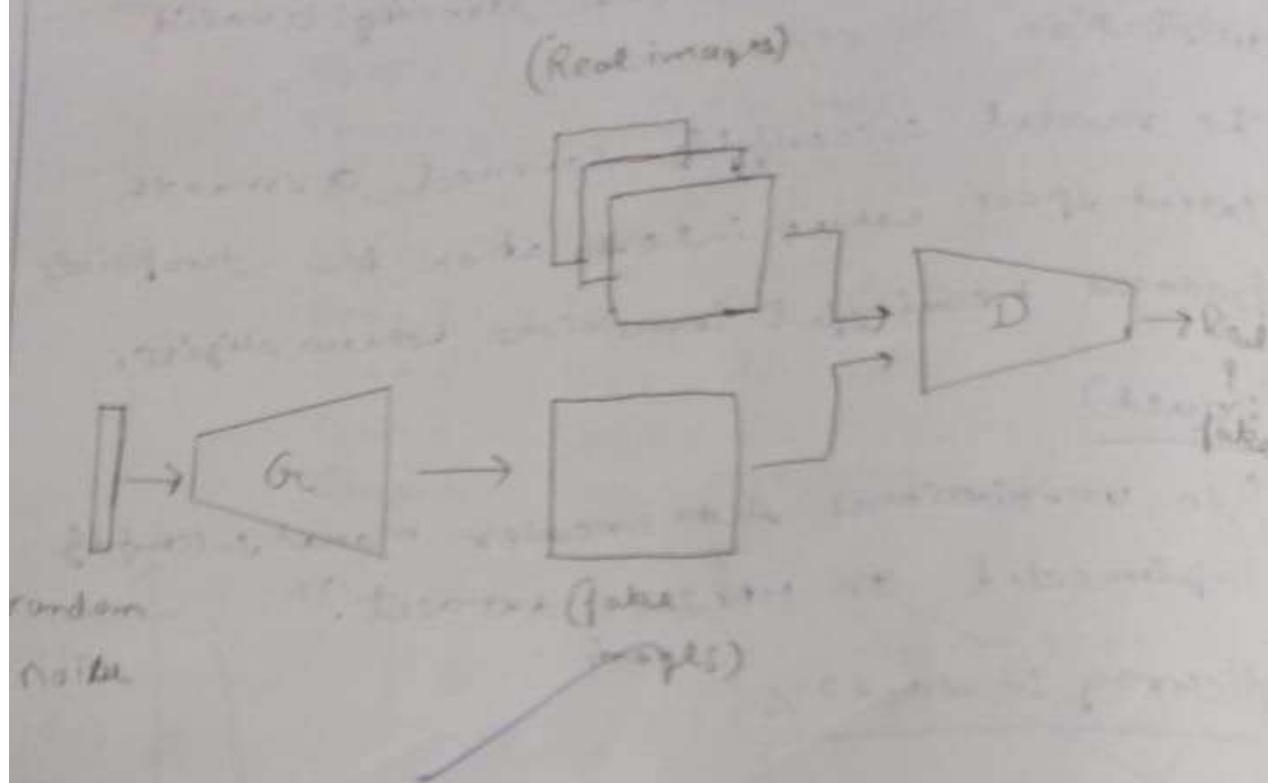
25-10-25

Lab 12:- (Implement a Deep Convolutional GAN to generate Complex Color Images)

Aim:- To implement a Deep Convolutional Generative Adversarial Network (DCGAN) that can generate complex color images.

Pseudo code:-

- Import libraries:-
torch, matplotlib.
- Load dataset:-
Use ~~labeled~~ ^{labeled} dataset
Normalize images.
- Define Generator Network:
- Define Discriminator Network
Input; Real or fake images.
- Initialize both networks and optimizers.
→ Adam optimizer.
- Training loop.
- Visualize.
 - Display generated color images.
 - Compare evolution across epochs.



Observation)

- During the initial epochs, generated images are random noise with no structure.
- As training progress, the generator learns color patterns, textures, and shapes resembling real images.
- The discriminator loss oscillates b/w 0 and 1.
- After sufficient training, DCGAN produces visually realistic and colorful synthetic images.
- The quality of generated images depends on dataset complexity, network depth and training stability.

Result)

((successfully implemented a DGAN capable of generating complex, realistic color images .

Ques

Loss

Output:- Epoch 1:- D: 0.6353
G: 1.4945

Epoch 2:- D: 0.2506
G: 2.1320

Epoch 3:- D: 0.1758
G: 2.0664

Epoch 4:- D: 0.10366
G: 2.1848

Epoch 5:- D: 0.0211
G: 0.1800

Epoch 6:- D: 0.0188
G: 0.4111

Epoch 7:- D: 0.0099
G: 0.1699

Epoch 8:- D: 0.2179
G: 2.5012

Epoch 9:- D: 0.0610
G: 0.37526

Epoch 10:- D: 1.3041
G: 1.1274

25-10-25

Lab13:- (Understanding the Architecture of a Pre-trained Model)

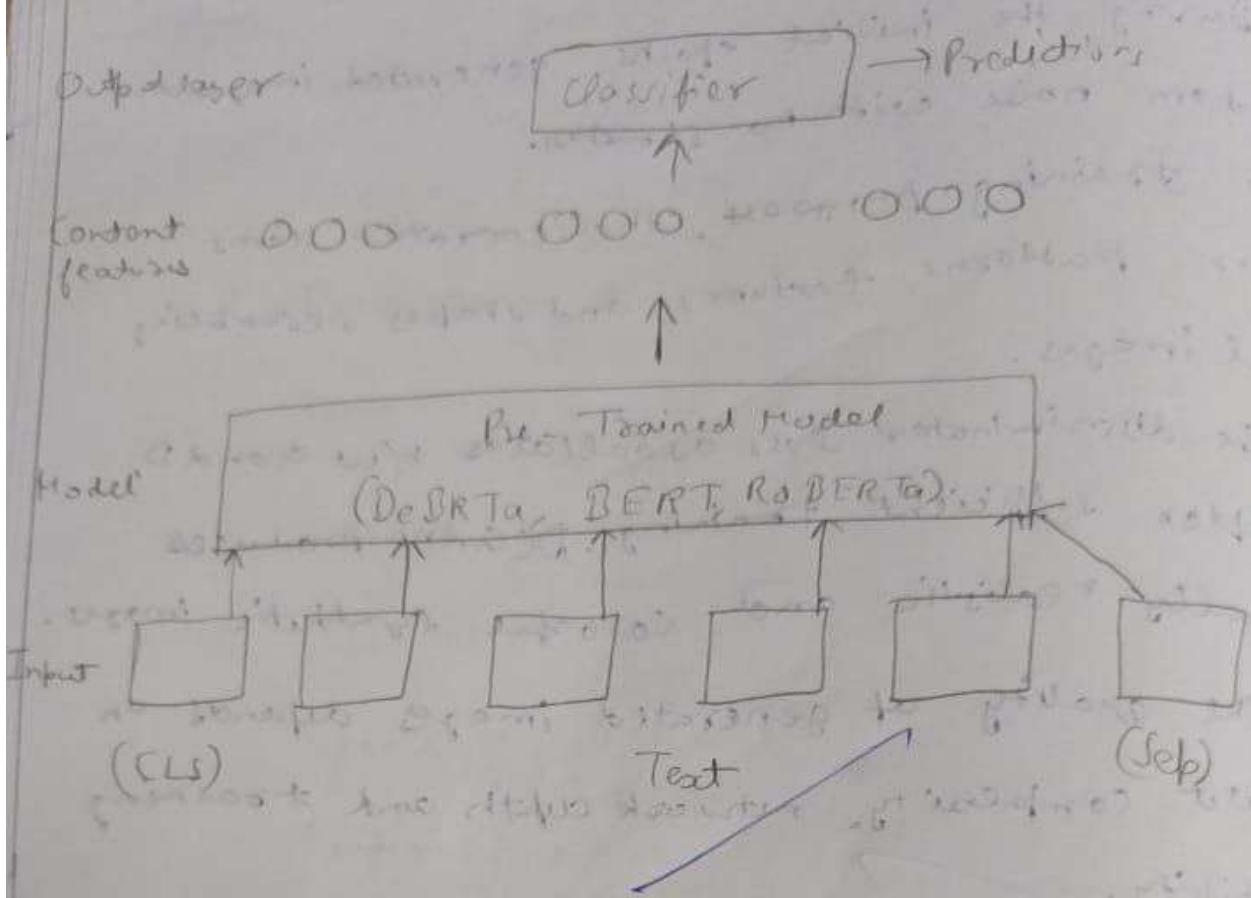
Aim:- To understand and analyze the architecture of a pre-trained deep learning model.

Pseudo code:- . Import req. libraries.

- Load a pre-trained model from pytorch (torchvision).
- Display the full architecture of the model.
- Count total trainable and non-trainable parameters
- Visualize layer types (Conv Pooling).
- Optionally pass a sample image through the model to verify dimensions.
- Analyze layer-by-layer flow and parameter size.

(Observation)

- The VGG16 model consists of 13 convolutional layers, 3 fully connected layers, and uses ReLU after each convolution.



The model ends with a softmax classifier.

The feature extractor part includes multiple Conv + Max Pool blocks, which progressively reduce spatial dimensions.

Total parameters are around 13.8 million.

Pre-trained weights help in transfer learning -

(contd.)

"The architecture and structure of the pre-trained model were successfully analyzed."

Qn

(Output)

	Top1Accuracy	Top5Accuracy	Parameters
VGG16	79.0%	94.5%	22.9M
Training	validation	Testing	
86.62%	91.05%	89.07%	

The screenshot shows a Jupyter Notebook interface with a dark theme. The top bar includes the file name "Untitled10.ipynb", a star icon, and a share icon. The menu bar has options: File, Edit, View, Insert, Runtime, Tools, Help. Below the menu is a toolbar with search, command, code, text, and run all buttons. On the right, there are icons for message, settings, and user profile, along with a "Connect 14" button.

The main area contains a code cell with the following Python script:

```
# PART 1: Import and Configuration
# -----
#
import os
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms, utils
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt

# Device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Hyperparameters
image_size = 32
nc = 3 # RGB
nz = 100 # latent vector size
ngf = 64 # generator feature maps
ndf = 64 # discriminator feature maps
batch_size = 128
num_epochs = 10
lr = 0.0002
beta1 = 0.5

OUT_DIR = "./dcgan_outputs"
os.makedirs(OUT_DIR, exist_ok=True)
```

At the bottom left, there are buttons for Variables and Terminal. A blue circular button with a white upward-pointing arrow is located at the bottom center.

The screenshot shows a Jupyter Notebook interface with a dark theme. The top bar includes the file name "Untitled10.ipynb", a search icon, and a refresh icon. The menu bar contains File, Edit, View, Insert, Runtime, Tools, and Help. The toolbar below the menu has icons for Commands, Code, Text, and Run all. The left sidebar features a tree view with "All" selected, showing various notebooks and files. The main area displays two code cells:

```
# -----
# PART 2: Dataset (CIFAR-10)
# -----  
  
transform = transforms.Compose([
    transforms.Resize(image_size),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5),
                      (0.5, 0.5, 0.5)), # normalize to [-1,1]
])  
  
dataset = datasets.CIFAR10(root='./data', download=True,
                           transform=transform)  
  
dataloader = DataLoader(dataset, batch_size=batch_size,
                           shuffle=True, num_workers=2)  
  
# -----
# PART 3: Generator & Discriminator
# -----  
  
# Weight initialization
def weights_init(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1:
        nn.init.normal_(m.weight.data, 0.0, 0.02)
    elif classname.find('BatchNorm') != -1:
        nn.init.normal_(m.weight.data, 1.0, 0.02)
        nn.init.constant_(m.bias.data, 0)  
  
# Generator
```

At the bottom, there are buttons for Variables and Terminal, and a blue run button.

Untitled10.ipynb ⭐ ⓘ

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect 14 ⓘ

```
# -----
# PART 4: Initialization
# 

netG = Generator(nz, ngf, nc).to(device)
netG.apply(weights_init)

netD = Discriminator(nc, ndf).to(device)
netD.apply(weights_init)

criterion = nn.BCELoss()
optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))

fixed_noise = torch.randn(64, nz, 1, 1, device=device)
real_label = 1.
fake_label = 0.

[1]: print("Training DCGAN for 10 epochs...")
for epoch in range(1, num_epochs + 1):
    for i, (data, _) in enumerate(dataloader, 0):
        # --- Train Discriminator ---
        netD.zero_grad()
        real = data.to(device)
        b_size = real.size(0)
        label = torch.full((b_size,), real_label, dtype=torch.float, device=device)
        output = netD(real)
        lossD_real = criterion(output, label)
        lossD_real.backward()
        D_x = output.mean().item()

        [2]: Variables Terminal
```

Untitled10.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect 14

```
output = netD(fake)
lossG = criterion(output, label)
lossG.backward()
D_G_z2 = output.mean().item()
optimizerG.step()

if i % 200 == 0:
    print(f"[{epoch}/{num_epochs}] [{i}/{len(dataloader)}] "
          f"Loss_D: {lossD.item():.4f} | Loss_G: {lossG.item():.4f}")
```

Training DCGAN for 10 epochs...

```
[1/10] [0/391] Loss_D: 1.4107 | Loss_G: 0.7604
[1/10] [200/391] Loss_D: 0.6353 | Loss_G: 1.4345
[2/10] [0/391] Loss_D: 0.2630 | Loss_G: 2.3100
[2/10] [200/391] Loss_D: 0.3006 | Loss_G: 2.2320
[3/10] [0/391] Loss_D: 0.1905 | Loss_G: 3.1562
[3/10] [200/391] Loss_D: 0.2758 | Loss_G: 2.0664
[4/10] [0/391] Loss_D: 0.0755 | Loss_G: 3.0738
[4/10] [200/391] Loss_D: 0.0566 | Loss_G: 3.1548
[5/10] [0/391] Loss_D: 0.0464 | Loss_G: 3.5390
[5/10] [200/391] Loss_D: 0.0258 | Loss_G: 3.8800
[6/10] [0/391] Loss_D: 0.0195 | Loss_G: 4.1255
[6/10] [200/391] Loss_D: 0.0158 | Loss_G: 4.4111
[7/10] [0/391] Loss_D: 0.0149 | Loss_G: 4.6282
[7/10] [200/391] Loss_D: 0.0099 | Loss_G: 4.8699
[8/10] [0/391] Loss_D: 0.0081 | Loss_G: 5.0301
[8/10] [200/391] Loss_D: 0.2179 | Loss_G: 2.5012
[9/10] [0/391] Loss_D: 0.0809 | Loss_G: 3.2313
[9/10] [200/391] Loss_D: 0.0610 | Loss_G: 3.7586
[10/10] [0/391] Loss_D: 0.2054 | Loss_G: 2.7939
[10/10] [200/391] Loss_D: 1.5041 | Loss_G: 1.1274
```

Untitled11.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect T4

```
import torch
import torch.nn as nn
import torchvision.models as models
from collections import OrderedDict
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

```
import torch
from torchvision import models
from torchvision.models import VGG16_Weights, ResNet50_Weights

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model_name = "vgg16" # or "resnet50"

if model_name.lower() == "vgg16":
    model = models.vgg16(weights=VGG16_Weights.DEFAULT).to(device)
elif model_name.lower() == "resnet50":
    model = models.resnet50(weights=ResNet50_Weights.DEFAULT).to(device)
else:
    raise ValueError("Supported: 'vgg16' or 'resnet50'")

model.eval()
print(f"\nLoaded model: {model_name.upper()} on {device}")

Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.cache/torch/hub/checkpoints/vgg16-397923af.pth
100%|██████████| 528M/528M [00:03<00:00, 146MB/s]
```

Variables Terminal

```
i 1 import torch
from torchvision import models
from torchvision.models import VGG16_Weights, ResNet50_Weights

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model_name = "resnet50" # or "resnet50"

if model_name.lower() == "vgg16":
    model = models.vgg16(weights=VGG16_Weights.DEFAULT).to(device)
elif model_name.lower() == "resnet50":
    model = models.resnet50(weights=ResNet50_Weights.DEFAULT).to(device)
else:
    raise ValueError("Supported: 'vgg16' or 'resnet50'")

model.eval()
print(f"\n\x1b[32m Loaded model: {model_name.upper()} on {device}\x1b[0m")

→ Downloading: "https://download.pytorch.org/models/resnet50-11ad3fa6.pth" to /root/.cache/torch/hub/checkpoints/resnet50-11ad3fa6.pth
100%|██████████| 97.8M/97.8M [00:00<00:00, 144MB/s]

\x1b[32m Loaded model: RESNET50 on cuda
```